


# Introduction to Leonardo

Author: CINECA Production Team  
[superc@cineca.it](mailto:superc@cineca.it)

Presenter: Gabriella Bettonte  
[g.bettonte@cineca.it](mailto:g.bettonte@cineca.it)

20th April 2024



## 2023 OVERVIEW

# HPC SYSTEMS

CINECA enables world-class scientific research by operating and supporting leading-edge supercomputing technologies and by managing a state-of-the-art and effective environment for the different scientific communities.

# CINECA




### MARCONI | 2017

3188 nodes  
48 cores per node  
612 TB RAM  
10 PFlops



### LEONARDO | 2023

4992 nodes  
**Booster Module:**  
32 core per node  
4 GPU NVIDIA Ampere custom  
**Data Centric Module:**  
56 cores per node  
SOON IN PRODUCTION  
110 PB Storage  
250 PFlops



### DGX | 2021

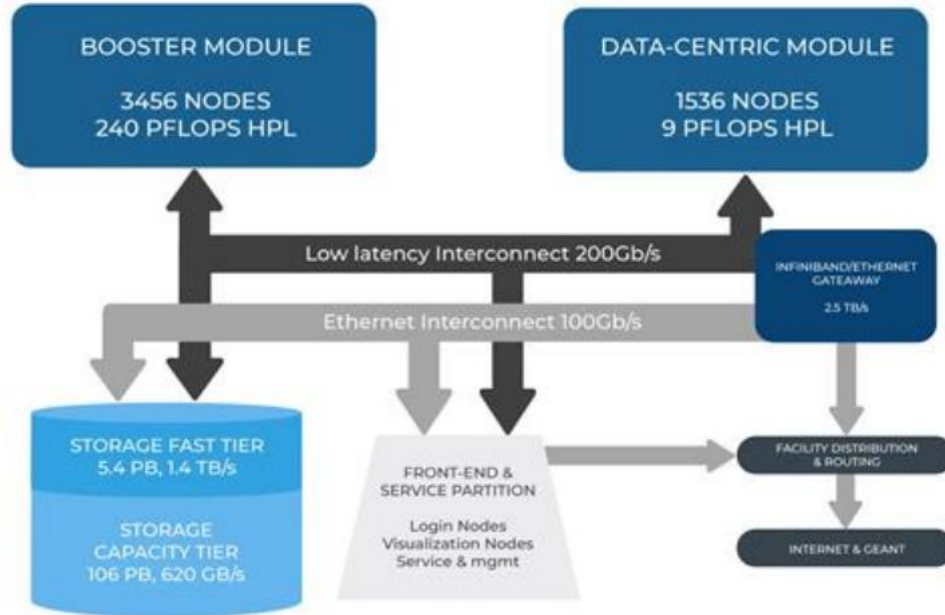
3 nodes  
128 cores per node  
8 GPU NVIDIA A100 per node  
100 TB Storage  
15 PFlops



### GALILEO100 | 2021

564 nodes  
48 cores per node  
2 GPU NVIDIA V100 per node  
~22 PB Storage  
2 PFlops

# Leonardo infrastructure and login nodes



## Atos BullSequana X430-E6

- Processors: **2 x CPU Intel Whitley ICP06, 32 cores Intel Ice Lake, 2.4 GHz**
- Hyper Threading is enabled!
- RAM: 512 (16x32) GB RAM DDR4 3200MHz
- 6TB disk in RAID1 configuration
- **NO GPUs**

# Booster (GPU) module

## Atos BullSequana X2135 “Da Vinci” blade


- 3456 nodes (1 node per blade)
- Processors: **1 x CPU Intel Xeon 8358,**  
**32 cores Intel Ice Lake, 2.6 GHz (ONE SOCKET!)**
- RAM: 512 (8 x 64) GB DDR4 3200 MHz
- Accelerators: **4 x NVidia custom Ampere GPU A100 SXM4 64 GB,**  
**NVLink 3.0**
- Internal network: NVIDIA Mellanox HDR DragonFly+ 200Gb/s
- **DISKLESS!!!**
- Shared (via infiniband) storage space: 106 PB Capacity tier storage  
+ 5.4 PB Fast tier storage



**Peak performance per node: about 89,4 TFlops**

**Peak performance: about 309 PFlops**

# Data Centric & General Purpose (CPU) module



**Peak performance per node: about 8.46 TFlops**  
**Peak performance: about 13 PFlops**

## BullSequana X2140 three-node CPU Blade

- 1536 nodes (512 blades)
- Processors: **2 x CPU Intel Xeon 8480+, 56 cores Intel Sapphire Rapids, 2.0 GHz**
- RAM: 256 (16x16) GB DDR5 4800MHz  
512 (16 x 32) GB DDR5 4800 MHz
- Infiniband: 1 x NVIDIA HDR cards 100 Gbps via PCIe Gen 5
- Disk: 1 x M.2 SSD 3,84 TB
- It will be in production soon.

# Visualization (viz) module

## 16 Atos BullSequana X450-E6

- Processors: 2 x CPU Intel Whitley ICP06,  
32 cores Intel Ice Lake, 2.4 GHz
- RAM: 512 GB (16x32) GB 3200MHz DDR4
- 2x Nvidia Quadro RTX8000 48GB PCIe
- 1x 6.4TB NVMe disks
- 1 x NIC HDR100 via Connect-X 6 card
- 2x Nic Connect-X 5 Ethernet 50 Gb/s
- **Not yet in production:** these nodes will be probably dedicated to the remote visualization or to the interactive computing service



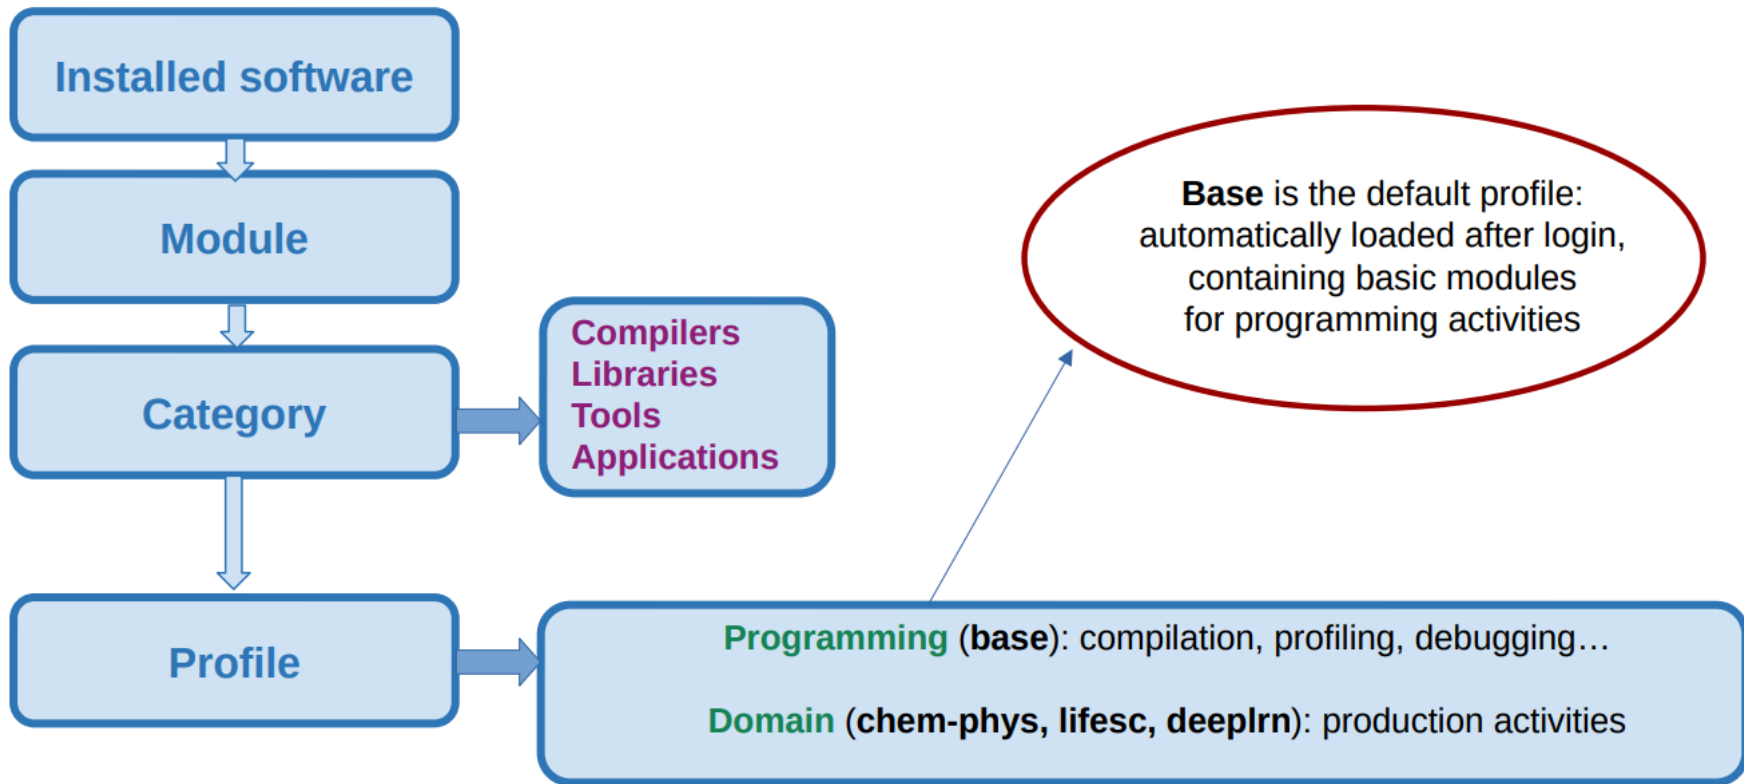
# Access to Leonardo

ssh [username@login.leonardo.cineca.it](ssh://username@login.leonardo.cineca.it)

For g100: ssh [username@login.g100.cineca.it](ssh://username@login.g100.cineca.it)

# Module environment

Any available software is offered on Leonardo in a **module environment**.  
The modules are organized in functional **categories** and collected in different **profiles**.





# Module environment

## \$ module av

```
----- /leonardo/prod/opt/modulefiles/profiles -----
profile/archive profile/base profile/candidate profile/chem-phys profile/deeplrn profile/lifesc profile/meteo profile/spoke7
----- /leonardo/prod/opt/modulefiles/base/archive -----
fake/1.0
----- /leonardo/prod/opt/modulefiles/base/dependencies -----
libxc/6.2.2--gcc--11.3.0-cuda-11.8
----- /leonardo/prod/opt/modulefiles/base/libraries -----
adios/1.13.1--openmpi--4.1.4--gcc--11.3.0      fftw/3.3.10--openmpi--4.1.4--nvhpc--23.1      nccl/2.14.3-1--gcc--11.3.0-cuda-11.8      parallel-netcdf/1.12.3--openmpi--4.1.4--gcc--11.3.0
adios/1.13.1--openmpi--4.1.4--nvhpc--23.1      gdal/3.5.3--gcc--11.3.0                        netcdf-c/4.9.0--gcc--11.3.0                parallel-netcdf/1.12.3--openmpi--4.1.4--nvhpc--23.1
blitz/1.0.2--gcc--11.3.0                        gsl/2.7.1--gcc--11.3.0                        netcdf-c/4.9.0--openmpi--4.1.4--gcc--11.3.0  parmetis/4.0.3--openmpi--4.1.4--gcc--11.3.0
boost/1.80.0--gcc--11.3.0                      hdf5/1.12.2--gcc--11.3.0                      netcdf-c/4.9.0--openmpi--4.1.4--nvhpc--23.1  parmetis/4.0.3--openmpi--4.1.4--nvhpc--23.1
boost/1.80.0--openmpi--4.1.4--gcc--11.3.0      hdf5/1.12.2--gcc--11.3.0-threadsafe            netcdf-fortran/4.6.0--gcc--11.3.0           petsc/3.18.1--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
boost/1.80.0--openmpi--4.1.4--nvhpc--23.1      hdf5/1.12.2--openmpi--4.1.4--gcc--11.3.0      netcdf-fortran/4.6.0--openmpi--4.1.4--gcc--11.3.0  petsc/3.18.1--openmpi--4.1.4--nvhpc--23.1-complex
cgal/5.4.1--gcc--11.3.0                        hdf5/1.12.2--openmpi--4.1.4--nvhpc--23.1      netcdf-fortran/4.6.0--openmpi--4.1.4--nvhpc--23.1  petsc/3.18.1--openmpi--4.1.4--nvhpc--23.1-cuda-11.8
cgal/5.4.1--openmpi--4.1.4--gcc--11.3.0        intel-oneapi-mkl/2022.2.1                     netlib-scalapack/2.2.0--openmpi--4.1.4--gcc--11.3.0  petsc/3.19.0--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
cineca-hpyc/2023.05                          intel-oneapi-mpi/2021.7.1                     netlib-scalapack/2.2.0--openmpi--4.1.4--nvhpc--23.1  proj/8.2.1--gcc--11.3.0
cudnn/8.4.0.27-11.6--gcc--11.3.0              intel-oneapi-tbb/2021.7.1                     openblas/0.3.21--gcc--11.3.0                rapids/2023.09
cutensor/1.5.0.3--gcc--11.3.0                 libmatio/1.1.11--gcc--11.3.0                  openblas/0.3.21--nvhpc--23.1                slate/2022.07.00--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
elpa/2021.11.001--openmpi--4.1.4--gcc--11.3.0-cuda-11.8  libzip/2.1.1--gcc--11.3.0                    openmpi/4.1.4--gcc--11.3.0-cuda-11.8         zlib/1.2.13--gcc--11.3.0
fftw/3.3.10--gcc--11.3.0                      magma/2.6.2--gcc--11.3.0-cuda-11.8             openmpi/4.1.4--nvhpc--23.1-cuda-11.8
fftw/3.3.10--openmpi--4.1.4--gcc--11.3.0      metis/5.1.0--gcc--11.3.0
----- /leonardo/prod/opt/modulefiles/base/tools -----
anaconda3/2022.05  emacs/28.2      gnuplot/5.4.3--gcc--11.3.0  nco/5.0.1--openmpi--4.1.4--gcc--11.3.0  singularity/3.8.7  texinfo/6.5
anaconda3/2023.03  git-lfs/3.1.2  intel-oneapi-vtune/2022.4.1  ncview/2.1.8--openmpi--4.1.4--gcc--11.3.0  snakemake/6.15.1  texinfo/6.5--gcc--11.3.0
cmake/3.24.3       git/2.38.1     jube/2.4.3                 ninja/1.11.1  spack/0.19.1-d71  valgrind/3.19.0--openmpi--4.1.4--gcc--11.3.0
curl/7.79.0       git/2.38.1--nvhpc--23.1  maven/3.8.4                openjdk/11.0.17_8  superc/2.0
----- /leonardo/prod/opt/modulefiles/base/compilers -----
cuda/11.8  intel-oneapi-compilers/2023.0.0  nvhpc/22.3  nvhpc/23.5  perl/5.36.0--gcc--11.3.0  python/3.10.8--gcc--8.5.0
gcc/11.3.0  llvm/15.0.4--gcc--11.3.0-cuda-11.8  nvhpc/23.1  perl/5.36.0--gcc--8.5.0  perl/5.36.0--nvhpc--23.1  python/3.10.8--gcc--11.3.0
```

# Module environment

Profile base → Automatically loaded

How to **load** additional profiles/modules?

```
$ module load profile/chem-phys
```

→ You can use tab for auto completion

Which modules did I **already load**?

```
$ module list
```

Currently Loaded Modulefiles:

1) profile/base 2) profile/chem-phys

*profile/base cannot be unloaded*

How to **unload** a profile/module?

*Specific profile/module*

```
$ module unload profile/chem-phys
```

*Unload all of them*

```
$ module purge
```

# Module environment

profile/base + profile/chem-phys → \$ module avail

----- /leonardo/prod/opt/modulefiles/profiles -----  
profile/archive profile/base profile/candidate profile/chem-phys profile/deeplrn profile/lifesc profile/meteo profile/spoke7

----- /leonardo/prod/opt/modulefiles/chem-phys/applications -----

amber/2022	kokkos/3.7.00--openmpi--4.1.4--gcc--11.3.0-cuda-11.8	quantum-espresso/7.2--openmpi--4.1.4--nvhpc--23.1-openblas-cuda-11.8
ams/2023.101	lammps/20220623--openmpi--4.1.4--gcc--11.3.0-cuda-11.8	thermoPW/1.8.1
cp2k/2023.2--openmpi--4.1.4--gcc--11.3.0-omp-cuda-11.8	namd/2.14--gcc--11.3.0-cuda-11.8	vasp/6.4.0
cp2k/master--openmpi--4.1.4--gcc--11.3.0-omp-cuda-11.8	nwchem/7.0.2--openmpi--4.1.4--gcc--11.3.0	yambo/5.1.1--openmpi--4.1.4--nvhpc--23.1
<b>gromacs/2021.7--openmpi--4.1.4--gcc--11.3.0-cuda-11.8</b>	plumed/2.8.1--openmpi--4.1.4--gcc--11.3.0	yambo/5.1.2--openmpi--4.1.4--nvhpc--23.1
<b>gromacs/2022.3--gcc--11.3.0-cuda-11.8</b>	plumed/2.9.0--openmpi--4.1.4--gcc--11.3.0	yambo/5.2.0--openmpi--4.1.4--nvhpc--23.1
<b>gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8</b>	quantum-espresso/7.2--openmpi--4.1.4--gcc--11.3.0-openblas	

**\$ module load gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8**

Autoload is not necessary → Different from other Cineca clusters

Loading **gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8**

Loading requirement: zlib/1.2.13--gcc--11.3.0 openmpi/4.1.4--gcc--11.3.0-cuda-11.8 fftw/3.3.10--openmpi--4.1.4--gcc--11.3.0

openblas/0.3.21--gcc--11.3.0 cblas/2015-06-06--gcc--11.3.0 gsl/2.7.1--gcc--11.3.0 plumed/2.8.1--openmpi--4.1.4--gcc--11.3.0



# Module environment

Loading a module

→ Define or modify the environment variables, allowing to use the executable or libraries

```
$ module load gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

**\$ module show <module\_name>/<version>** → Prints information about the module: dependencies, paths

```
[otrocon1@login05 ~]$ module show gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

```
/leonardo/prod/opt/modulefiles/chem-phys/applications/gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8:
```

```
module-whatis {GROMACS is a molecular dynamics package primarily designed for simulations of proteins, lipids and nucleic acids. It was originally developed in the Biophysical Chemistry department of University of Groningen, and is now maintained by contributors in universities and research centers across the world.}
```

```
module load fftw/3.3.10--openmpi--4.1.4--gcc--11.3.0
```

```
module load openblas/0.3.21--gcc--11.3.0
```

```
module load openmpi/4.1.4--gcc--11.3.0-cuda-11.8
```

```
module load plumed/2.8.1--openmpi--4.1.4--gcc--11.3.0
```

```
conflict gromacs
```

```
prepend-path GROMACS_LIB /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/lib64
```

```
prepend-path LIBRARY_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/lib64
```

```
prepend-path LD_LIBRARY_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/lib64
```

```
prepend-path GROMACS_INC /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/include
```

```
prepend-path GROMACS_INCLUDE /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/include
```

```
prepend-path C_INCLUDE_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/include
```

```
prepend-path CPLUS_INCLUDE_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/include
```

```
prepend-path CPATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/include
```

```
prepend-path PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/bin
```

```
prepend-path MANPATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/share/man
```

```
prepend-path PKG_CONFIG_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/lib64/pkgconfig
```

```
prepend-path CMAKE_PREFIX_PATH /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex/.
```

```
setenv GROMACS_HOME /leonardo/prod/spack/03/install/0.19/linux-rhel8-icelake/gcc-11.3.0/gromacs-2022.3-owzxiwojzrytaodpf2r7dvd63jflbvex
```

# Module environment

```
$ module load profile/lifesc
```

```
$ module load gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

```
$ module help gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

---

Module Specific Help for `/leonardo/prod/opt/modulefiles/chem-phys/applications/gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8:`

modulefile "gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8"

using help from `/cineca/prod/opt/helps/gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8`

Example of batch script for MPI+CUDA version with or without plumed:

1) Hybrid MPI/OpenMP job on 2 nodes without plumed:

```
#!/bin/bash
#SBATCH --job-name job_name
#SBATCH -N2 --ntasks-per-node=4
#SBATCH --cpus-per-task=8
#SBATCH --time=24:00:00
#SBATCH --account=<account_nr>
#SBATCH --partition=boost_usr_prod
#SBATCH --gres=gpu:4
```

```
module load profile/lifesc
module load gromacs/2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

```
export OMP_NUM_THREADS=8
```

```
cmd="gmx_mpi mdrun -s topol.tpr -deffnm md -ntomp 8 -v -nb gpu -pme gpu -npme 1 -pin on -nstlist 500"
mpirun -np 8 $cmd
```

*The script example will be different in other cluster(s)*



# Module environment

How to find a module that I do not know in which profile is it?

`$ modmap -m <module_name>` → a command that looks for a module in all profiles

```
[otrocon1@login05 ~]$ modmap -m gromacs
Profile: archive
Profile: base
Profile: chem-phys
        applications
            gromacs
                2021.7--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
                2022.3--gcc--11.3.0-cuda-11.8
                2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
Profile: deeplrn
Profile: lifesc
        applications
            gromacs
                2021.7--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
                2022.3--gcc--11.3.0-cuda-11.8
                2022.3--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
Profile: meteo
Profile: spoke7
        applications
            gromacs
                2021.7--openmpi--4.1.4--gcc--11.3.0-cuda-11.8
```

module load profile/quantum

```
----- /leonardo/prod/opt/modulefiles/quantum/applications -----  
ocean/6.6.0 qiskit/0.44.3 qmatcha_tea/1.1.4
```

ml qmatcha\_tea/

```
[gbettont@login02 ~]$ ml qmatcha_tea/  
Loading qmatcha_tea/1.1.4  
Loading requirement: gmp/6.2.1 mpfr/4.1.0 mpc/1.2.1 gcc/11.3.0 python/3.10.8--gcc--8.5.0 cutensor/1.5.0.3--gcc--11.3.0 cuda/11.8  
zlib/1.2.13--gcc--11.3.0 openmpi/4.1.4--gcc--11.3.0-cuda-11.8 openblas/0.3.21--gcc--11.3.0  
netlib-scalapack/2.2.0--openmpi--4.1.4--gcc--11.3.0
```

**Please remember that on G100 you need also this command --> module load autoloader**

# Production environment: how to submit my simulations

You have compiled your code! Congratulations!!!  
...now, what to do with it?



The **production environment** takes care of all the aspects related to the actual execution of the program you want to run. Time to “produce” some results!!



# Your work areas

## (recap)

### **\$HOME:**

Permanent, backed-up, and local to LEONARDO. 50 Gb of quota. For source code or important input files.

### **\$SCRATCH:**

Large, parallel filesystem. No quota. Run your simulations and calculations here. A cleaning policy will delete all your files older than 40 days (not active yet).

### **\$WORK:**

Similar to \$SCRATCH, but the content is shared among all the users of the same account.  
1 Tb of quota (no cleaning policy).

### **\$PUBLIC:**

A small area to share installations, permanent but not backed-up. 50 Gb of quota.

use the command **cindata** to get info on your disk occupation

<https://wiki.u-gov.it/confluence/display/SCAIUS/UG2.5%3A+Data+storage+and+FileSystems>

# Login nodes

If you have a serial program, the most intuitive thing to do is to just launch `./myprogram` wherever you are.

When you log into LEONARDO, you find yourself in one of the four login nodes, selected in round robin fashion to balance out the load of users.

Interactive runs on login nodes are strongly discouraged and should be limited to **short test runs**

**There are per user limits** on cpu-time (10 minutes)

**IMPORTANT:** avoid running large parallel applications on the front-ends!!

LEONARDO is a general purpose system used by hundreds of users.

# Compute nodes: jobs & scheduler

What we actually want most of the time is to gain access to the compute nodes to exploit their power

Like in any HPC cluster, LEONARDO allows you to run your simulations by submitting “jobs” to the compute nodes

Your job is then taken in consideration by a scheduler, that adds it to a queuing line and allows its execution when the resources required are available



The operative  
scheduler on LEONARDO is **SLURM**

SLURM stands for "Simple Linux Utility for Resource Management"

- Allocating access to resources
- Job starting, executing and monitoring
- Queue of pending jobs management

# Compute nodes: jobs & scheduler

The scheme for a SLURM job script is as follows:

**#!/bin/bash**

**#SLURM directives**

**variables environment**

**execution line**

# Jobscript example

```
#!/bin/bash
#SBATCH -t 1:00:00
#SBATCH -N 2
#SBATCH --ntasks-per-node=16
#SBATCH --cpus-per-task=2
#SBATCH --gres=gpu:4
#SBATCH --mem=10GB
#SBATCH -o job.out
#SBATCH -e job.err
#SBATCH -p boost_usr_prod
#SBATCH -A <my_account>
module load openmpi
export OMP_PROC_BIND=true
mpirun -n 32 ./myprogram
```

# Slurm directives

**#SBATCH --job-name=myname, -J myname**

Defines the name of your job

**#SBATCH --output=job.out, -o job.out**

Specifies the file where the standard output is directed (default=slurm-<Pid>)

**#SBATCH --error=job.err, -e job.err**

Specifies the file where the standard error is directed (default=slurm-<Pid>)

# Slurm directives: resource requirements

**#SBATCH --nodes=1, -N 1**

**#SBATCH --ntasks-per-node=8**

**#SBATCH --cpus-per-task=4**

**#SBATCH --gres=gpu:4**

**#SBATCH --mem=10000**      # mem=0 equals to full memory

**nodes** – number of compute nodes

**ntasks-per-node** – number of tasks per node (max. 32)

**cpus-per-task** – number of cpus to be assigned to each task

$\text{ntasks-per-node} * \text{cpus-per-task} \leq 32$

**gres=gpu:x** – number of GPUs for each node (x=1..4)

**mem** – memory allocated for each node (max=494000 MB).

# Slurm directives: walltime and partitions

**#SBATCH --time=00:30:00, -t 00:30:00**

Specifies the maximum duration of the job. The maximum time allowed depends on the partition used

**Pro-tip:** the less walltime you ask, the faster your job will enter in execution. Think about it!

**#SBATCH --partition=boost\_usr\_prod, -p boost\_usr\_prod**

**#SBATCH --qos=boost\_qos\_dbg, -q boost\_qos\_dbg (optional)**

Specifies the “partition”, a.k.a. the specific set of nodes among which your job can search for resources. Optionally you can specify a QoS (Quality of Service) for jobs with particular purposes, like debugging or large production



# Available partitions and QoS on LEONARDO

SLURM partition	Job QOS	# cores/# GPU per job	max walltime	max running jobs per user/ max n. of nodes/cores/GPUs per user	priority	notes
lrd_all_serial (default)	<i>normal</i>	max = 4 physical cores (8 logical cpus)  max mem = 30800 MB	04:00:00	1 node / 4 cores / 30800 MB	40	No GPUs Hyperthreading x2
boost_usr_prod	<i>normal</i>	max = 32 nodes	24:00:00		40	
	boost_qos_dbg	max = 2 nodes	00:30:00	2 nodes / 64 cores / 8 GPUs	80	
	boost_qos_bprod	min = 33 nodes  max = 256 nodes	24:00:00	256 nodes	60	runs on 1536 nodes min is 33 FULL nodes
	boost_qos_lprod	max = 3 nodes	4-00:00:00	3 nodes / 12 GPUs	40	

## Notes:

- the partition **lrd\_all\_serial** run on front-end nodes, and as such it is not subject to accounting and can be used for free
- to use the QoS **boost\_qos\_bprod**, the minimum requirement is for cpus, gpus, nodes and mem to be over the regular limit...that's why the nodes have to be asked in full!

# Slurm directives: accounting

**#SBATCH --account=<my\_account>, -A <my\_account>**

Specifies the account to use the CPU hours from.

As an user, you have access to a limited number of CPU hours to spend. They are not assigned to users, but to **projects** and are shared between the users who are working on the same project (i.e. your research partners). Such projects are called **accounts** and are a different concept from your username.

You can check the status of your account with the command “*saldo -b*”, which tells you how many CPU hours you have already consumed for each account you’re assigned at (a more detailed report is provided by “*saldo -r*”).

```
[amarani@login01 ~]$ saldo -b
```

account	start	end	total (local h)	localCluster Consumed(local h)	totConsumed (local h)	totConsumed %	monthTotal (local h)	monthConsumed (local h)
cin_staff	20110323	20300323	200000002	16382068	50270876	25.1	864553	785338
cin_propro	20220427	20301231	500000	2568	2695	0.5	4731	0
cin_saldo	20230524	20300323	10	0	0	0.0	0	0
cin_sudo	20230524	20300323	10	0	0	0.0	0	0
FUSIO_TEST_4	20161116	20231231	1000	191	191	19.2	11	149
cin_external_6	20150319	20231231	20000	7695	7695	38.5	186	103
cin_extern01 3	20170210	20231231	5000	0	0	0.0	59	0

# Submitting a job

You have created your jobscript! Congratulations!!!  
...now, what to do with it?



## **sbatch**

`sbatch <job_script>`

Your job will be submitted to the SLURM scheduler and executed when there will be nodes available (according to your priority and the partition you requested)

## **squeue -u**

`squeue -u <username>`

Shows the list of all your scheduled jobs, along with their status (idle, running, closing, ...)  
It also shows you the job id required for other SLURM commands

# Other Slurm commands - 1

```
[amarani0@login02 r206n06]$ sbatch submit_gpu.sh
Submitted batch job 382861
[amarani0@login02 r206n06]$ squeue -u amaranio
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
382861	m100_usr_	submit_g	amarani0	PD	0:00	1	(Resources)

## scontrol show job

scontrol show job <job\_id>

Provides a long list of informations for the job requested.

In particular, if your job isn't running yet, you'll be notified about the reason it is not starting and, if it is scheduled with top priority, you will get an estimated start time

## scancel

scancel <job\_id>

Removes the job (queued or running) from the scheduled job list by killing it

# Other Slurm commands - 2

## sinfo

sinfo -p <partition name>

sinfo -l

sinfo -N -l -p m100\_usr\_prod

Provides information about SLURM nodes and partitions

## sacct

sacct OPTIONS <job\_id>

displays accounting data for all jobs and job steps in the SLURM job accounting log or Slurm database. Can also see the jobs already completed or cancelled

```
[amarani0@login02 r206n06]$ sacct -X --format=jobid,user,account,start,nnodes,elapsed
```

JobID	User	Account	Start	NNodes	Elapsed
382860	amarani0	cin_staff	2020-06-26T12:04:51	1	00:03:01
382861	amarani0	cin_staff	2020-06-26T12:07:52	1	00:03:01

```
[amarani0@login02 r206n06]$
```

# Documentation

Our userguide goes into more depth about all the aspects described during this presentation. In particular, we suggest:

<https://wiki.u-gov.it/confluence/display/SCAIUS/UG3.2%3A+LEONARDO+UserGuide#UG3.2:LEONARDOUserGuide-Productionenvironment>

Production environment on LEONARDO

<https://wiki.u-gov.it/confluence/display/SCAIUS/UG2.5%3A+Data+storage+and+FileSystems>

Work areas and filesystem

<https://wiki.u-gov.it/confluence/display/SCAIUS/UG2.4%3A+Accounting>

Accounting and budget consumption

<https://wiki.u-gov.it/confluence/display/SCAIUS/UG2.6.1%3A+How+to+submit+the+job+-+Batch+Scheduler+SLURM>

Batch scheduler Slurm

*coming soon...*

Analysis on thread and task affinity on LEONARDO